

In [6]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing, svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression
8 from sklearn.preprocessing import StandardScaler
```

In [7]:

```
1 df=pd.read_csv(r"C:\Users\Welcome\Documents\bottle1.csv")
2 df
```

C:\Users\Welcome\AppData\Local\Temp\ipykernel_6724\3464836710.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.

```
df=pd.read_csv(r"C:\Users\Welcome\Documents\bottle1.csv")
```

Out[7]:

| | Cst_Cnt | Btl_Cnt | Sta_ID | Depth_ID | Depthm | T_degC | Salnty | O2ml_L | STheta | O2Sa |
|--------|---------|---------|----------------|--|--------|--------|---------|--------|----------|-------|
| 0 | 1 | 1 | 054.0 056.0 | 19- 4903CR- HY-060- 0930- 05400560- 0000A-3 | 0 | 10.500 | 33.4400 | NaN | 25.64900 | Na |
| 1 | 1 | 2 | 054.0 056.0 | 19- 4903CR- HY-060- 0930- 05400560- 0008A-3 | 8 | 10.460 | 33.4400 | NaN | 25.65600 | Na |
| 2 | 1 | 3 | 054.0 056.0 | 19- 4903CR- HY-060- 0930- 05400560- 0010A-7 | 10 | 10.460 | 33.4370 | NaN | 25.65400 | Na |
| 3 | 1 | 4 | 054.0 056.0 | 19- 4903CR- HY-060- 0930- 05400560- 0019A-3 | 19 | 10.450 | 33.4200 | NaN | 25.64300 | Na |
| 4 | 1 | 5 | 054.0 056.0 | 19- 4903CR- HY-060- 0930- 05400560- 0020A-7 | 20 | 10.450 | 33.4210 | NaN | 25.64300 | Na |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | . |
| 864858 | 34404 | 864859 | 093.4 026.4 | 20- 1611SR- MX-310- 2239- 09340264- 0000A-7 | 0 | 18.744 | 33.4083 | 5.805 | 23.87055 | 108.7 |
| 864859 | 34404 | 864860 | 093.4 026.4 | 20- 1611SR- MX-310- 2239- 09340264- 0002A-3 | 2 | 18.744 | 33.4083 | 5.805 | 23.87072 | 108.7 |
| 864860 | 34404 | 864861 | 093.4 026.4 | 20- 1611SR- MX-310- 2239- 09340264- 0005A-3 | 5 | 18.692 | 33.4150 | 5.796 | 23.88911 | 108.4 |
| 864861 | 34404 | 864862 | 093.4 026.4 | 20- 1611SR- MX-310- 2239- 09340264- 0010A-3 | 10 | 18.161 | 33.4062 | 5.816 | 24.01426 | 107.7 |

| | Cst_Cnt | Btl_Cnt | Sta_ID | Depth_ID | Depthm | T_degC | Salnty | O2ml_L | STheta | O2Sat |
|--------|---------|---------|--------|--|--------|--------|---------|--------|----------|-------|
| | | | | 20-1611SR-093.4MX-310-026.42239-09340264-0015A-3 | | | | | | |
| 864862 | 34404 | 864863 | | | 15 | 17.533 | 33.3880 | 5.774 | 24.15297 | 105.6 |

864863 rows × 74 columns

```
In [11]: 1 df=df[['Salnty','T_degC']]
          2 df.columns=['sal','temp']
```

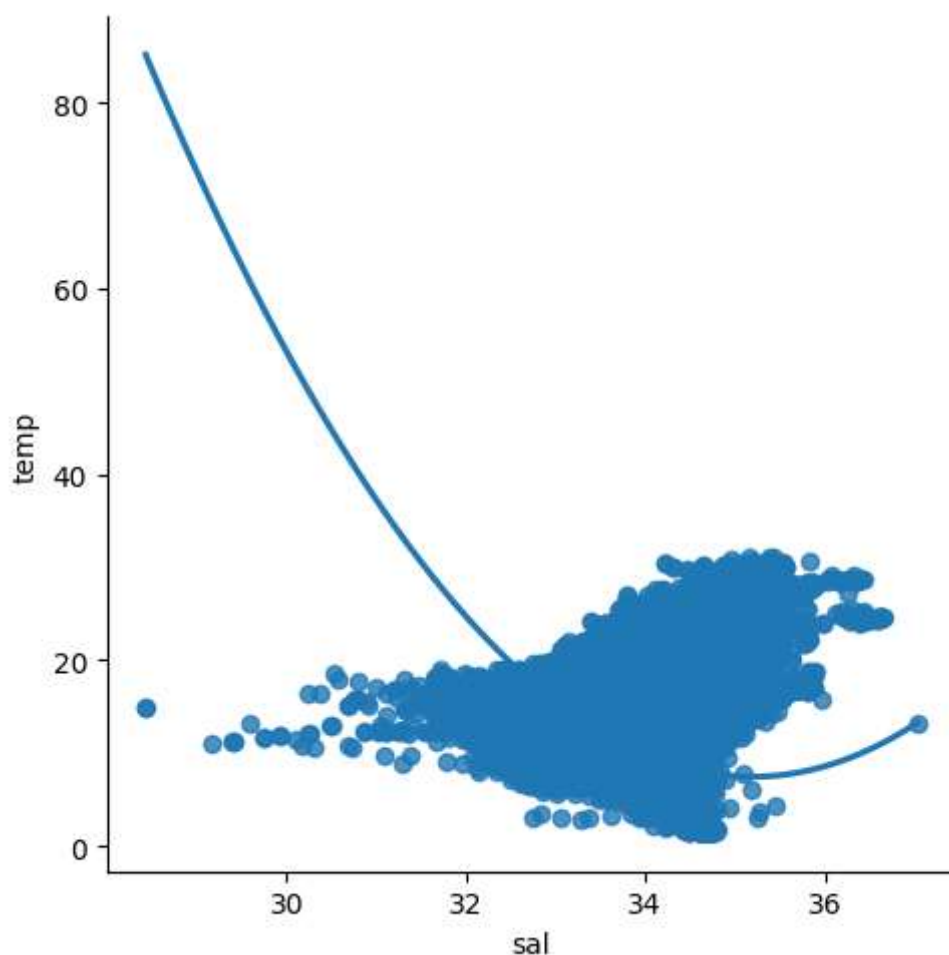
```
In [12]: 1 df.head(10)
```

Out[12]:

| | sal | temp |
|---|--------|-------|
| 0 | 33.440 | 10.50 |
| 1 | 33.440 | 10.46 |
| 2 | 33.437 | 10.46 |
| 3 | 33.420 | 10.45 |
| 4 | 33.421 | 10.45 |
| 5 | 33.431 | 10.45 |
| 6 | 33.440 | 10.45 |
| 7 | 33.424 | 10.24 |
| 8 | 33.420 | 10.06 |
| 9 | 33.494 | 9.86 |

```
In [13]: 1 sns.lmplot(x='sal',y='temp',data=df,order=2,ci=None)
```

```
Out[13]: <seaborn.axisgrid.FacetGrid at 0x176ab76f010>
```



```
In [14]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 864863 entries, 0 to 864862  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0   sal      817509 non-null   float64  
1   temp     853900 non-null   float64  
dtypes: float64(2)  
memory usage: 13.2 MB
```

In [15]: 1 df.describe()

Out[15]:

| | sal | temp |
|-------|---------------|---------------|
| count | 817509.000000 | 853900.000000 |
| mean | 33.840350 | 10.799677 |
| std | 0.461843 | 4.243825 |
| min | 28.431000 | 1.440000 |
| 25% | 33.488000 | 7.680000 |
| 50% | 33.863000 | 10.060000 |
| 75% | 34.196900 | 13.880000 |
| max | 37.034000 | 31.140000 |

In [16]: 1 df.fillna(method='ffill')

Out[16]:

| | sal | temp |
|--------|---------|--------|
| 0 | 33.4400 | 10.500 |
| 1 | 33.4400 | 10.460 |
| 2 | 33.4370 | 10.460 |
| 3 | 33.4200 | 10.450 |
| 4 | 33.4210 | 10.450 |
| ... | ... | ... |
| 864858 | 33.4083 | 18.744 |
| 864859 | 33.4083 | 18.744 |
| 864860 | 33.4150 | 18.692 |
| 864861 | 33.4062 | 18.161 |
| 864862 | 33.3880 | 17.533 |

864863 rows × 2 columns

In [17]: 1 df.fillna(value=0,inplace=True)

In [18]: 1 df.isnull().sum()

Out[18]: sal 0
temp 0
dtype: int64

In [19]: 1 x=np.array(df['sal']).reshape(-1,1)
2 y=np.array(df['temp']).reshape(-1,1)

```
In [20]: 1 df.isna().any()
```


```
Out[20]: sal      False  
temp      False  
dtype: bool
```

```
In [21]: 1 df.dropna(inplace=True)
```

```
In [22]: 1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
2 reg=LinearRegression()  
3 reg.fit(X_train,y_train)  
4 print(reg.score(X_test,y_test))
```


```
0.014953114878132445
```

```
In [23]: 1 y_pred=reg.predict(X_test)  
2 plt.scatter(X_test,y_test,color='b')  
3 plt.scatter(X_test,y_pred,color='k')  
4 plt.show()
```




```
In [24]: 1 df500=df[:][:500]  
2 sns.lmplot(x="sal",y="temp",data=df500,order=1,ci=None)
```

Out[24]: <seaborn.axisgrid.FacetGrid at 0x176ab76f040>




```
In [25]: 1 df500.fillna(method='ffill',inplace=True)
2 X=np.array(df500['sal']).reshape(-1,1)
3 y=np.array(df500['temp']).reshape(-1,1)
4 df500.dropna(inplace=True)
5 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
6 reg=LinearRegression()
7 reg.fit(X_train,y_train)
8 print("Regression:",reg.score(X_test,y_test))
9 y_pred=reg.predict(X_test)
10 plt.scatter(X_test,y_test,color="b")
11 plt.plot(X_test,y_pred,color='k')
12 plt.show()
```

Regression: 0.042321355512622616



```
In [26]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model=LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred=model.predict(X_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2 score:",r2)
```

R2 score: 0.042321355512622616


conclusion:Linear Regression is best fit for the model

Ridge and Lasso Regression

```
In [27]: 1 from sklearn.linear_model import Ridge
          2 from sklearn.linear_model import RidgeCV
          3 from sklearn.linear_model import Lasso
```

```
In [28]: 1 plt.figure(figsize = (10, 10))
          2 sns.heatmap(df.corr(), annot = True)
```

Out[28]: <Axes: >



```
In [29]: 1 features = df.columns[0:2]
2 target = df.columns[-1]
3 #X and y values
4 X = df[features].values
5 y = df[target].values
6 #split
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, r
8 print("The dimension of X_train is {}".format(X_train.shape))
9 print("The dimension of X_test is {}".format(X_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)
```

The dimension of X_train is (605404, 2)

The dimension of X_test is (259459, 2)

```
In [30]: 1 lr = LinearRegression()
2 #Fit model
3 lr.fit(X_train, y_train)
4 #predict
5 #prediction = lr.predict(X_test)
6 #actual
7 actual = y_test
8 train_score_lr = lr.score(X_train, y_train)
9 test_score_lr = lr.score(X_test, y_test)
10 print("\nLinear Regression Model:\n")
11 print("The train score for lr model is {}".format(train_score_lr))
12 print("The test score for lr model is {}".format(test_score_lr))
13
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0


```
In [31]: 1 ridgeReg = Ridge(alpha=10)
2 ridgeReg.fit(X_train,y_train)
3 #train and test scorefor ridge regression
4 train_score_ridge = ridgeReg.score(X_train, y_train)
5 test_score_ridge = ridgeReg.score(X_test, y_test)
6 print("\nRidge Model:\n")
7 print("The train score for ridge model is {}".format(train_score_ridge))
8 print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.999999999723243

The test score for ridge model is 0.9999999997231402

```
In [32]: 1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
3 plt.plot(lr.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color
4 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersiz
5 plt.xticks(rotation = 90)
6 plt.legend()
7 plt.show()
```



```
In [33]: 1 print("\nLasso Model: \n")
2 lasso = Lasso(alpha = 10)
3 lasso.fit(X_train,y_train)
4 train_score_ls =lasso.score(X_train,y_train)
5 test_score_ls =lasso.score(X_test,y_test)
6 print("The train score for ls model is {}".format(train_score_ls))
7 print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.0

The test score for ls model is -1.9031696447013857e-05

```
In [34]: 1 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind =
2
```

Out[34]: <Axes: >




```
In [35]: 1 #Using the Linear CV model
2 from sklearn.linear_model import LassoCV
3 #Lasso Cross validation
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state
5 #score
6 print(lasso_cv.score(X_train, y_train))
7 print(lasso_cv.score(X_test, y_test))
```

0.9999999994806811

0.9999999994806712

```
In [36]: 1 #plot size
2 plt.figure(figsize = (10, 10))
3 #add plot for ridge regression
4 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
5 #add plot for lasso regression
6 plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6
7 #add plot for linear model
8 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersiz
9 #rotate axis
10 plt.xticks(rotation = 90)
11 plt.legend()
12 plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
13 plt.show()
```



```
In [37]: 1 #Using the Linear CV model
2 from sklearn.linear_model import RidgeCV
3 #Ridge Cross validation
4 ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train,
5 #score
6 print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
7 print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.9999999986797505

The train score for ridge model is 0.9999999986778121

```
In [38]: 1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(X,y)
4 print(regr.coef_)
5 print(regr.intercept_)
6
```

[0. 0.94934511]

0.5401219631067828